

Journal: Ecological Applications

Manuscript type: Method

**Smart environmental DNA sampling strategies: a spatial autocorrelation method
and interactive planner for survey design**

Running headline: Efficient eDNA survey design

Gledis Guri^{1*}, Ryan P. Kelly¹, Andrew O. Shelton², Owen R. Liu³, Krista M. Nichols²

¹School of Marine and Environmental Affairs, University of Washington, Seattle, Washington,
USA

²Conservation Biology Division, Northwest Fisheries Science Center, NOAA Fisheries, 2725
Montlake Blvd E, Seattle, Washington 98112, USA

³School of Aquatic and Fishery Sciences, University of Washington, Seattle, Washington, USA

*Corresponding author. Email: gguri@uw.edu

Open Research statement: Data are archived in Zenodo (Guri, 2026;
<https://doi.org/10.5281/zenodo.21141367>). Novel code (R and Stan scripts that reproduce all
analyses) and the eDNA Survey Planner are provided for peer review in a public GitHub
repository (<https://github.com/gledguri/Smart-eDNA-sampling>), and the planner can be used
directly in a web browser at <https://gledguri.github.io/Smart-eDNA-sampling/>. Upon acceptance,
the version of record of the code and the planner will be permanently archived in Zenodo.

Key words: biodiversity monitoring; decision-support tool; environmental DNA; fisheries
management; Gaussian process; metabarcoding; sampling design; spatial autocorrelation; survey
cost-effectiveness

Abstract

Environmental DNA (eDNA) surveys are increasingly used to monitor marine biodiversity, yet managers have little quantitative guidance on how many samples to collect, and because species differ in patchiness, no single sample size serves all. We present a method that converts a target prediction error and a prior estimate of a taxon's spatial autocorrelation length-scale (ρ) into a minimum sampling density. Bayesian Gaussian process models are fitted to pilot eDNA data, observed and simulated fields are progressively thinned to quantify how reduced effort degrades prediction and parameter recovery, and the results are summarized in an invertible error–effort equation implemented in a free, browser-based survey planner. In a case study combining eDNA data for 12 fish species (358 samples from five depth strata along the U.S. West Coast) with 200 simulated fields ($\rho = 25\text{--}1000$ km), marginal gains became negligible at 4–5 samples per 10,000 km² per depth layer for smoothly distributed species ($\rho > 500$ km), whereas patchily distributed species ($\rho < 50$ km) required 8–11, over twice the effort for equal precision. Depth-specific means, which underpin abundance indices, were recovered reliably under sparse sampling. The planner lets survey designers set defensible, species-specific effort before a cruise.

Introduction

Biodiversity monitoring programmes underpin decisions about where to fish, what to protect and how to report progress against national and international conservation targets, yet the budgets available to them are effectively fixed. How many samples to collect, and where, is therefore a management decision as much as a statistical one (Yoccoz et al., 2001). Survey design must balance the need for sufficient data for detecting trends and changes in population

while also minimizing cost (i.e., maximizing sampling efficiency) (Stein & Ettema, 2003; Xu et al., 2015; Zhao et al., 2022).

The emergence of environmental DNA (eDNA) sampling has revolutionized monitoring capabilities, offering a non-invasive method to detect species presence and estimate relative abundance (Gold et al., 2021; Shelton et al., 2022; Taberlet et al., 2012). Recent synthesis work shows that eDNA methods have matured sufficiently for key uncertainties to be quantified and incorporated into existing fisheries stock-assessment frameworks, enabling the development of cost-effective abundance indices that complement conventional surveys (Baetscher et al., 2025; Guri et al., 2024; Kirtane et al., 2021; Salter et al., 2026). Because agencies are now evaluating eDNA as an operational component of biodiversity assessment rather than as an exploratory tool, explicit and defensible guidance on how much to sample has become a prerequisite for uptake rather than a refinement of it.

While increased sampling intensity generally reduces variance in population estimates and improves the precision of spatial pattern detection in species distribution models (Schweiger et al., 2016), this relationship exhibits diminishing returns where beyond a certain threshold, additional sampling effort produces marginal gains in statistical precision that may not justify the associated costs and logistical demands (Perret et al., 2022; Zhao et al., 2022).

We conceptualize species spatial distributions as realizations of a Gaussian Process (GP), in which spatial dependence is characterized by a species-specific length-scale parameter (ρ) describing the distance over which observations remain correlated. We hypothesize that the sampling density required to characterize a species' distribution depends on this spatial length-scale. A species with a long length-scale varies gradually across space and can therefore be reconstructed from fewer samples, whereas a species with a short length-scale exhibits more

rapid spatial variation and consequently requires denser sampling (Figure 1). Thus, knowledge of species-specific spatial autocorrelation can provide a quantitative basis for informing sampling effort (Fortin et al., 1989). Yet there is still no general, transferable procedure for deciding how many eDNA samples a survey needs, and design guidance is rarely used unless it is accessible to those who plan surveys. Here we present such a procedure together with a free, interactive tool that implements it. Using paired simulated and empirical eDNA data for species that differ in their spatial distributions, the method quantifies how spatial autocorrelation and sampling effort affect (i) the predictive accuracy of species distribution models and (ii) the accuracy of GP parameter estimation, and summarizes the result in an equation that can be solved for the sampling effort required. Our objective is deliberately operational: to provide survey planners with guidance that can be applied before a cruise, based on a prior estimate of the spatial autocorrelation of target taxa and the level of precision required. We describe the method and the planner, demonstrate both with a multispecies eDNA survey of the California Current, and outline applications to survey planning, monitoring and management.

Description of the method

The method links sampling effort to the accuracy of spatial inference within a common GP framework and proceeds in four steps (Figure 2): (1) fit a GP model to existing or pilot eDNA data to characterize taxon-specific spatial structure; (2) progressively thin the observed data, refit the model and quantify how prediction accuracy degrades; (3) repeat the thinning on simulated fields spanning a wide range of spatial autocorrelation, at resolutions no field campaign can achieve, to extend inference beyond the observed effort and evaluate parameter recovery against known values; and (4) summarize prediction error as a function of sampling density and spatial autocorrelation in a regression that can be inverted to return the sampling

density required for a user-defined error tolerance. Step 4 is implemented in the eDNA Survey Planner, so that end users need only a prior estimate of spatial autocorrelation and a target error. We describe each step generically, giving the values used in the case study where relevant.

Modelling framework: Gaussian processes

The GP approach is grounded in the assumption that species distributions and abundances change in a smooth and continuous manner across space (or other covariates), reflecting correlated responses to spatial, environmental, and biological covariates (Lewy & Kristensen, 2009). At the core, in simple terms, GPs posit that observations made at proximate locations or covariate values should be similar (Dormann et al., 2007). We assert that many species can be well described using relatively simple covariance models (Thorson et al., 2015). Specifically, we use an isotropic, squared exponential covariance model for the distribution of a species. Let X $\{x_1, x_2, \dots, x_N\}$ be the set of N observed spatial locations. Then species (eDNA) abundance or density Y across all sampled locations (X) is modeled as,

$$Y \sim MVN(\mu, \mathbf{K} + \sigma^2 \mathbf{I}), \#1$$

where μ is the global mean abundance, σ^2 is the non-spatial variance, $\mathbf{I}$ is an $n \times n$ identity matrix and $\mathbf{K}(X, X)$ is the covariance matrix whose $(i, j)^{\text{th}}$ entry is given by the squared exponential kernel function,

$$\mathbf{K}(X, X) = k(x_i, x_j) = \alpha^2 \exp\left(-\frac{(x_i - x_j)^2}{2\rho^2}\right) \#2$$

where α^2 is the amplitude (i.e., the marginal variance of the spatial process, controlling the overall magnitude of variation in abundance), ρ is the length-scale parameter (in $x_i - x_j$ units; hereafter spatial autocorrelation).

Step 1: Estimating taxon-specific spatial structure

Given log eDNA concentration values (Y) paired with spatial coordinates (X) for each sampling location (x ; $N = 358$ in the case study), we used the GP model (Eq. 1–2) to infer the underlying spatial structure and learn key parameters governing the spatial field, including the depth-specific mean (μ_d), spatial autocorrelation (ρ), amplitude (α^2), and non-spatial variance (σ^2). Additionally, we used the model to predict eDNA concentrations (Z) across a spatial prediction grid (Q) using Eq 3. In the case study, we defined this grid (Q) as 18,916 discrete point estimates, representing the centroids of approximately 8×8 km spatial cells distributed across five depth layers. The number of points differed among depth layers (4,000, 4,000, 4,000, 3,622, and 3,294 points, respectively) because deeper layers covered a smaller spatial extent than shallower layers due to bathymetry.

For a given species, we compute the covariance matrix among observed locations, $\mathbf{K}(X, X')$, using the squared exponential kernel and the non-spatial variance term ($\sigma^2 \mathbf{I}$). Using the Cholesky decomposition we solved for $(\mathbf{K}(X, X') + \sigma^2 \mathbf{I})^{-1}(Y - \mu_d)$ and we combined it with the cross-covariance between observed and prediction points, $\mathbf{K}(X, Q)$, to yield the predictions at each given spatial point (Q):

$$Z = \mu_d + \mathbf{K}(X, Q)^\top (\mathbf{K}(X, X') + \sigma^2 \mathbf{I})^{-1} (Y - \mu_d) \quad \#3$$

Step 2: Quantifying information loss under reduced sampling effort

To evaluate how sampling effort (E) influences spatial reconstruction and prediction accuracy, we generated a series of reduced-effort observation datasets (Y_E coupled with X_E) by randomly thinning the observation set (Y) and its spatial covariate (X) to fixed sample sizes (in the case study, $N = 330, 300, 260, 220$, and 150 ; Figure 2B). Each reduced-effort observation dataset (Y_E) together with its covariate (X_E) was refit with the same GP model (Eq 1-2) and produced effort-specific spatial predictions (S_E) on the full grid Q , by substituting S_E for Z , Y_E for

Y and X_E for X in Eq. 3. This framework enables direct, like-for-like comparisons of reconstructed fields across sampling intensities since the predicted field is the same (Q).

Sampling effort (E) was quantified as sampling density:

$$E = \frac{N}{A \times H} \times 10000 \#4$$

where A is the spatial area of the study domain (200,000 km² in the case study), and H is the number of strata, here depth layers (d ; $H = 5$ in the case study), yielding E in units of samples per 10,000 km² per depth layer.

For each model prediction, we calculated the prediction error (ε_{1E}), defined as the root mean squared error between spatial fields predicted from the full dataset (Z) and those predicted from each reduced sampling scenario (S_E) as follows (Figure 2B):

$$\varepsilon_{1E} = \sqrt{\frac{1}{n} \sum_{i=1}^n (Z_i - S_{iE})^2} \#5$$

with n being the total number of Q grid points ($n = 18,916$) where the predictions (Z and S_E) are projected. We repeated this thinning-fitting-prediction process across 20 random realizations for each thinning intensity.

Step 3: Extending the range of sampling effort with simulated fields

Because no observed dataset can span the full range of sampling effort of interest ($N = 358$ in the case study), the third step extends the effort range by simulation. We generated *in-silico* eDNA concentration data using the GP model (Eq. 1-2) across the same spatial domain and prediction grid (Q) used in the empirical analysis.

The simulations were designed to explore a wide range of spatial structures by varying the GP length-scale parameter (ρ) across ten levels: 25, 50, 75, 100, 200, 300, 400, 500, 800, and 1000 km (Appendix S1: Table S1). For each value of ρ we generated (arbitrarily) 20 independent

replicate spatial fields, resulting in a total of 200 simulated datasets (10 ρ levels $\times$ 20 replicates; Appendix S1: Table S1). Similar to GP inference of the empirical data, within each simulation, ρ was held constant across depth layers, but each layer had a separate mean (μ_d) drawn randomly from a normal distribution with mean of 0 and a standard deviation of 3. We focused only on spatial autocorrelation across latitude and longitude, allowing the eDNA concentration at each depth to be independent. The amplitude (α^2) and the non-spatial variance (σ^2) in Eq. 1 and 2 were fixed at 4 and 3, respectively, to mirror the empirical estimates, where α^2 ranged from 2.3 to 4.2 across species and σ^2 was approximately 3 (Figure 3C).

Given the large number of projected eDNA points (Q), fitting a GP model using all locations ($\approx 20,000$ input values) was computationally prohibitive. We therefore drew a subset of 4,000 points from each of the 200 simulated datasets as the “complete” observation dataset (Y) and its covariate (X). Following a procedure identical to that used for the empirical data, we first fit the GP model (Eq. 1–2) to each complete simulated dataset (Y) and covariate (X) to infer its governing parameters (μ_d , ρ , α^2 , and σ^2) and thereafter predict eDNA concentrations (Z) across the full spatial grid (Q) using Eq. 3. This prediction (Z) served as the “true” spatial field against which all subsequent reduced-effort predictions were compared (Figure 2A).

To evaluate how sampling effort influences prediction accuracy on the simulated data, we generated a series of reduced-effort observation datasets (Y_E) and covariates (X_E) by randomly thinning each complete set (Y) and its covariate (X) to progressively smaller sample sizes ($N = 2000, 1000, 600, 358, 330, 300, 260, 220$, and 150 ; Figure 2A). Following the same procedure used for the empirical data, we fit the GP model (Eq. 1-3) to estimate the governing parameters and generate an effort-specific spatial prediction (S_E) on the full spatial grid (Q). We then

quantified the prediction error (ε_{2E}) for each scenario by calculating the root mean squared error between the “true” field (Z) and the reduced-effort field (S_E), as defined in Eq. 5 (Figure 2A).

Additionally, to evaluate the influence of sampling effort on GP parameter estimation, we compared the estimated parameters ($\rho, \mu_d, \sigma^2, \alpha^2$) from effort-specific datasets (Y_E) and covariates (X_E) to the true values used to generate the simulated fields. Differences between estimated and simulated parameter values were quantified using the root mean squared error (hereafter τ).

Step 4: Solving for the sampling density required by a stated objective

To examine how sampling effort and spatial autocorrelation jointly shape prediction error, we fit a log-linear model to each dataset separately (empirical: ε_1 ; simulated: ε_2), isolating each factor's contribution and enabling comparison across datasets.

We modeled prediction error (ε) as a function of sampling density (E) and spatial autocorrelation (ρ), including their interaction using a linear model (*lm* function) in R as follows:

$$\log(\varepsilon) = \omega + \beta \log(E) + \theta \log(\rho) + \gamma \log(E) \log(\rho) \quad \#6$$

where ω is the intercept, β slope of sampling density (E), θ is the slope of spatial autocorrelation (ρ) and γ is the coefficient on the interaction between the two latter factors.

For a given ρ , Eq. 6 reduces to a power law in E that can be inverted to return the minimum sampling density (E^*) required to achieve a target prediction error (ε^*):

$$E^* = \exp\left(\frac{\log(\varepsilon^*) - \omega - \theta \log(\rho)}{\beta + \gamma \log(\rho)}\right) \quad \#7$$

The number of samples required per stratum of area A is $A E^*/10,000$, rounded up to a whole number, and the total across H strata is H times that number. Conversely, a fixed budget of N_h samples per stratum supports a survey area of $10,000 N_h/E^*$ at the target error. Eq. 7 has a solution only while the fitted error declines with effort ($\beta + \gamma \log(\rho) < 0$), which holds throughout the calibrated range ($\rho = 25\text{--}1000$ km; $E = 1.5\text{--}20$ samples per 10,000 km² per stratum).

Benchmarking simulated against observed fields

Because the empirical benchmark (ε_1) was defined by the 358 observed samples, whereas the simulation benchmark (ε_2) used 4,000 samples, the two errors are not directly comparable (Figure 2). We therefore repeated the thinning on each simulated field using its $N = 358$ subsample as the reference (ε_3), fitted Eq. 6 to ε_3 and to ε_1 , and compared the resulting coefficients to test whether the simulations reproduce the error structure of real eDNA fields (Appendix S1: Section S2).

All the GP models (equation 1 and 2) were fitted in a Bayesian framework using Stan probabilistic programming language (Stan Development Team, 2023), with weakly informative priors, and the model in Eq. 6 was fitted using the *lm()* function in R (R Core Team, 2024); sampler settings, convergence diagnostics and prior distributions are given in Appendix S1: Section S3 and Table S2.

Implementation: the eDNA Survey Planner

To make Step 4 usable without statistical software, we implemented Eq. 6–7 in the eDNA Survey Planner, a self-contained browser application (a single HTML file) that runs offline, requires no installation and uploads no data (Appendix S1: Figure S6; Open Research statement). It answers the two questions survey designers most often face. “How many samples?” takes a target error (ε^*), a prior estimate of ρ , the survey area per stratum and the number of strata, and returns the required sampling density, the samples per stratum and the total site–stratum observations. “How much area?” takes ε^* , ρ and a sample budget per stratum, and returns the area that budget can support at the target error.

Because ε is a root mean squared error on the natural-log scale, the planner translates it into a concentration range: around a reference concentration C , predictions typically fall within

exp($\log(C) \pm \varepsilon$), a multiplicative factor of e^ε (e.g., $\varepsilon = 0.8$ corresponds to about 45–223 copies/L around a true 100 copies/L; Appendix S1: Figure S3). An interactive error–effort curve shows where the target lies on the curve of diminishing returns, and a one-dimensional GP simulation displays a true field, the samples required to reach ε^* and the resulting reconstruction, alongside one quarter of that effort. Spatial autocorrelation can be entered directly or chosen from presets (25–1,000 km).

Default coefficients are those of Eq. 6 estimated from the simulations, but they can be edited, so that programs can substitute a calibration for their own region, taxa or assay by applying Steps 1–4 with the archived code. Results outside the calibrated range are calculated but flagged as extrapolations, and the planner states that ε is an average error across the survey domain rather than a pointwise confidence interval. Recommendations follow the central fitted curve; uncertainty in the coefficients is not yet propagated.

Case study: a multispecies eDNA survey of the California Current

The case study shows how each step is executed on real data and tests whether the simulation-based error model holds where the true field is unknown.

Study system and eDNA data

Environmental DNA samples were collected during the Joint U.S.-Canada Integrated Ecosystem and Pacific Hake Acoustic-Trawl Survey aboard the NOAA ship *Bell M. Shimada* in summer 2019, from central California to northern Washington (38.36°–48.46° N, 123.53°–125.93° W; study domain of approximately 200,000 km²). Seawater was collected with Niskin bottles at five target depths (0, 50, 150, 300 and 500 m) and filtered onboard, yielding 358 site–depth samples from 152 sampling locations. Species-specific eDNA concentrations (log copies/L) were estimated for 12 fish species by jointly modeling 12S rRNA metabarcoding reads

(MiFish primers), a mock community to correct for amplification bias and qPCR measurements of Pacific hake (Guri et al., 2024, 2025). The focal assemblage comprised pelagic, mesopelagic and demersal taxa characteristic of the California Current ecosystem: *Clupea pallasii*, *Engraulis mordax*, *Leuroglossus stilbius*, *Microstomus pacificus*, *Sardinops sagax*, *Scomber japonicus*, *Sebastes entomelas*, *Stenobranchius leucopsarus*, *Tactostoma macropus*, *Tarletonbeania crenularis*, *Thaleichthys pacificus* and *Trachurus symmetricus*. Sampling, laboratory and bioinformatic details are given in Appendix S1: Section S1.

Spatial structure of the focal species (Step 1)

Across species, depth-specific global mean eDNA concentrations generally declined with increasing sampling depth (Figure 3). Species-specific length-scale estimates varied considerably across species (Figure 3A), likely reflecting differences in habitat use and movement behavior. Surface-dwelling species tended toward shorter spatial autocorrelation lengths and showed more pronounced differences between surface and deep mean eDNA concentrations (μ ; Figure 3B), though exceptions were evident. Deep-dwelling species broadly exhibited longer spatial autocorrelations (Figure 3A) and relatively stable mean eDNA concentrations across depth strata, although this pattern was not universal. *Tactostoma macropus*, for instance, showed a comparatively short length-scale despite its deep-water affinity.

Information loss in the observed data (Step 2)

Prediction error in empirical data decreased with increasing sampling density across all species (Appendix S1: Figure S1 and S2). However, this reduction was strongly modulated by the degree of spatial autocorrelation (Appendix S1: Figure S2). Species characterized by short spatial autocorrelation (small ρ values) consistently exhibited higher prediction error (ε_I) across all sampling densities, whereas species with wider spatial autocorrelation structures (e.g., $\rho >$

500 km) showed substantially lower prediction errors (ε_1) even under comparatively sparse sampling regimes.

Information loss in simulated fields (Step 3)

Prediction error in simulation data (ε_2) decreased non-linearly with increasing sampling effort and with higher spatial autocorrelation (Figure 4A). These patterns were supported by a log-log linear model where the prediction error (ε) decreased exponentially for every exponential increase in sampling density ($\beta = -0.176$) and spatial autocorrelation ($\theta = -0.228$), with an additional negative interaction ($\gamma = -0.0684$) and with an intercept of $\omega = 1.65$. For instance, at a sampling intensity of 6 samples per 10,000 km², species with spatial autocorrelation values of 50 km had mean prediction errors of 1 (Figure 4A) indicating that on average every point predicted the eDNA concentration will differ from the predicted true concentration (Z) by a factor of 2.72 (for perspective a concentration of 100 copies/L will be predicted between 37-270 copies/L; Appendix S1: Figure S3). In contrast, for the same sampling effort (i.e., 6 samples per 10,000 km²) but species with spatial autocorrelation $\rho > 500$ km had errors below 0.4 indicating an average of over- and underestimation by a factor of 1.5 (for perspective a concentration of 100 copies/L will be predicted between 67-150 copies/L; Appendix S1: Figure S3). When quantifying the marginal reduction in prediction error ($\Delta\varepsilon_2$) across successive increases in sampling density (ΔE), improvements were greatest at low sampling effort and declined rapidly with additional sampling, with $\Delta\varepsilon_2$ approaching zero at higher sampling densities (Figure 4B). Marginal gains effectively plateaued when sampling effort increased above 10 samples per 10,000 km², depending on the ρ (Figure 4B). Importantly, diagnostic checks revealed no systematic bias in the predictions, so the observed variation reflects pure precision (Appendix S1: Figure S4). For interpretative purposes, we adopt an arbitrary threshold of $\Delta\varepsilon_2/\Delta E = -0.05$ to

denote negligible marginal gains. However, this value is arbitrary and may be redefined according to study-specific accuracy requirements.

GP parameter recovery (Step 3)

Sampling effort and spatial autocorrelation strongly influenced the ability of the GP model to recover the parameters used to generate the simulated spatial fields (Figure 5). Estimation error for the depth-specific mean (μ_d) decreased with increasing sampling density for simulated fields with small spatial autocorrelation but showed the opposite pattern for fields with large spatial ranges (Figure 5C). Specifically, τ (root mean squared error between estimated and simulated) for μ_d increased with sampling effort when spatial autocorrelation exceeded 400 km, indicating that mean estimates became increasingly biased in highly spatially structured fields. Additionally, the recovery of the spatial autocorrelation length-scale (ρ) was highly sensitive to sampling density (Figure 5A). At low sampling effort ($E < 4$), estimates of ρ showed substantial error, particularly for simulations with large spatial ranges. However, estimation accuracy improved rapidly as sampling effort increased, with τ declining sharply and stabilizing once sampling densities were higher than 6 samples per 10,000 km², beyond which additional sampling produced only marginal improvements in ρ estimation.

Estimation error for the amplitude parameter (α^2) declined with increasing sampling density for simulations with small spatial autocorrelation, while remaining relatively stable across sampling densities for simulations with larger spatial ranges (Figure 5D). Overall improvements in α^2 estimation were therefore limited once sampling effort increased above 6 samples per 10,000 km². The non-spatial variance parameter (σ^2) was consistently well recovered across sampling densities (Figure 5B). τ declined sharply with increasing sampling density and approached near-zero values once sampling densities exceeded 6 samples per 10,000 km².

Agreement between simulated and observed error models

Prediction error (ε) followed similar patterns using both empirical and simulation datasets. Specifically, the estimated intercept (ω), the influence of sampling effort (β), the effect of spatial autocorrelation (θ), and their interaction term (γ) all showed consistent directional effects and aligned closely with the 1:1 line, demonstrating a high degree of structural agreement between the empirical and simulated error models (Appendix S1: Figure S1).

Applying the planner (Step 4)

To illustrate Step 4, consider a survey of the case-study domain (200,000 km² per depth layer, five layers) that aims to predict eDNA concentrations to within a typical factor of about 2.2 of the reference field ($\varepsilon^* = 0.8$). For a patchy taxon ($\rho = 50$ km), Eq. 7 returns $E^* \approx 9.1$ samples per 10,000 km² per depth layer, or 183 samples per layer (915 site–depth samples in total); for a smoothly distributed taxon ($\rho = 500$ km), the same target requires $E^* \approx 2.1$, or 43 samples per layer (215 in total). The 2019 survey ($E \approx 3.6$) is therefore expected to meet this target for smooth taxa ($\varepsilon \approx 0.59$) but not for patchy ones ($\varepsilon \approx 1.2$). Run in reverse, its budget of about 72 samples per layer would achieve $\varepsilon^* = 0.8$ for a patchy taxon over approximately 79,000 km², about 40% of the domain. The planner performs these calculations interactively (Appendix S1: Figure S6).

Discussion

We present a method, and a planner that implements it, for translating species-specific spatial autocorrelation into the sampling effort required by eDNA spatial surveys. The case study shows that optimal sampling effort is not universal but depends critically on both the spatial autocorrelation of target species and the specific inferential objectives of the survey.

Diminishing returns of sampling effort on prediction accuracy

The value of additional sampling effort is not fixed but depends on the spatial structure of the target species (Xu et al., 2015; Zhao et al., 2022). Prediction error decreased non-linearly with sampling effort and with increasing spatial autocorrelation, confirming that homogeneously distributed species are inherently easier to reconstruct than patchier ones at equivalent effort (Wagner & Fortin, 2005). This interaction matters because it shows the value of additional samples is not universal, where the same increase in effort can produce substantial gains for one species and only marginal gains for another (García-Machado et al., 2024). Our results move beyond a general statement of sampling efficiency, providing a quantitative framework for understanding when additional effort is informative and when it yields diminishing returns.

Our adopted threshold for negligible marginal gains ($\Delta\epsilon_2 = -0.05$) serves as a practical heuristic, but acceptable error is inherently context dependent. For instance, studies requiring high-resolution quantitative estimates might justify high sampling effort to achieve stricter precision, while, broad-scale biodiversity screenings might tolerate looser precision, allowing for even sparser sampling (Perkins et al., 2021; Radinger et al., 2019). Moreover, the species spatial autocorrelation parameter plays an additional role in defining the decision on sampling effort. Highly autocorrelated species (e.g., $\rho > 500$ km) reach the precision threshold of negligible improvement at much lower sampling densities (roughly 4 to 5 samples per 10,000 km²), conversely, patchily distributed species ($\rho < 50$ km) require a denser sampling network (roughly 8 to 11 samples per 10,000 km²) before the marginal gains drop below the threshold. Once the required precision and ρ are known (from prior literature, pilot sampling or expert judgement), the fitted log–log model (Eq. 6; Figure 4B; Appendix S1: Figure S5) can be inverted (Eq. 7) to solve for the minimum sampling density E required to reach a target error, a calculation the planner performs interactively.

Applying our simulation-derived thresholds to our empirical data further illustrates this context dependency. Extrapolating the findings from the simulations, species (with high spatial autocorrelations) such as *Leuroglossus stilbius* and *Microstomus pacificus* would require only 4 samples per 10,000 km² to achieve a negligible marginal gain in precision (defined here arbitrary as $\Delta\epsilon = -0.05$). In contrast such as *Trachurus symmetricus*, *Thaleichthys pacificus*, *Stenobranchius leucopsarus*, and *Tarletonbeania crenularis* (species with low spatial autocorrelation) would require approximately 9 samples per 10,000 km² to reach the same relative threshold (Liu et al., 2026). The remaining species fall within the range of 4-9 samples per 10,000 km² (for context this is slightly higher than our current empirical grid and CalCOFI's standard 75-station grid with 3.6 and 3.4 samples per 10000 km² respectively), consistent with their intermediate spatial autocorrelation values. However, given that our analysis is based on 12 fish species, extending this conclusion to other taxa or ecosystems should be done with caution, and additional species may exhibit different scaling relationships.

Although we report effort per 10,000 km², the underlying point is simpler: what matters is how far apart samples are relative to a species' autocorrelation. A sampling effort E puts samples roughly $\sqrt{(10,000/E)}$ km apart, so the thresholds above can equivalently be stated as a ratio of sample spacing to length-scale: $d/\rho \approx 0.1$ for smoothly distributed species ($E \approx 4-5$ with $\rho \approx 500$ km), but $d/\rho \approx 0.6-0.7$ for patchily distributed ones ($E \approx 8-11$ with $\rho \approx 50$ km). The same reasoning should therefore carry to other systems and scales, though the ratio is not a single fixed number, because error depends on E and ρ jointly (Eq. 6).

Assumptions and limitations of the method

Extrapolating directly from simulations to natural systems requires care because real eDNA fields may violate the simplifying assumptions of the stationary, isotropic squared

exponential kernel. Spatial structure may vary across the domain (non-stationarity), differ by direction (anisotropy), or include sharp transitions (fronts, plume boundaries, shelf-break structure) that are not well represented by a globally smooth covariance function (Paciorek & Schervish, 2006). In addition, eDNA observation error is unlikely to be purely Gaussian, with variability arising from sources such as PCR bias, compositional constraints, incomplete calibration, and heterogeneity in filtration or extraction efficiency (Bessey et al., 2020; Gloor et al., 2017; Kelly et al., 2019; Shaffer et al., 2025). The modelling framework used here accounts for these sources (Guri et al., 2025), so they are unlikely to bias the results presented here. Finally, ecological processes such as migration, schooling, habitat discontinuities, or patchy prey fields can generate multi-scale patterns that depart from a single characteristic length-scale (Wagner & Fortin, 2005). These features may help explain elevated prediction error at intermediate autocorrelation levels in the empirical data relative to simulations, where additional sources of spatial complexity are not represented.

Furthermore, the simulation assumes a stationary GP with fixed amplitude and noise and a single length-scale across depths, whereas real eDNA fields may have variable variance structure, multiple characteristic scales, and dependence on environmental covariates such as temperature, salinity, chlorophyll, and bathymetry (Liu et al., 2026). Our prediction error (ϵ_2) is also computed against a "complete" prediction from a 4,000-point subset (Z) rather than the full grid, a limitation most visible in the secondary thinning comparison, where the reference field is constrained to $N = 358$ samples, making prediction error higher and more stochastic.

Intrinsically, eDNA integrates ecological information through space and across time via transport and decay (Brandão-Dias et al., 2023; Fukaya et al., 2021). This means that even if organisms are locally patchy, the observed eDNA field may be smoother than the organism field,

depending on hydrodynamic conditions and degradation rates. In our framework, such integration is embedded within the inferred length-scale, reinforcing that ρ should be interpreted as an effective length scale of the observed eDNA field rather than a purely ecological niche parameter (Jo et al., 2025). Integrating eDNA observations with independent abundance data (trawl, acoustics) would allow some degree of separation of the spatial autocorrelation parameter, enabling clearer ecological interpretation of ρ and improving the transferability of sampling-effort recommendations (Guri et al., 2024).

Spatial scale also sets an important boundary on the applicability of these results. The relationships identified here are expected to hold most strongly when interpreted relative to the extent and resolution of the study domain, rather than as fixed thresholds transferable across all systems. In a small study domain, the autocorrelation scale of the observed eDNA field may exceed the survey extent, causing the field to appear nearly uniform and making μ_d difficult to estimate independently of the spatial process (Paciorek, 2010; Redfern et al., 2006).

Perspectives and applications

The method turns a single estimable quantity, the spatial autocorrelation of the target taxon, into a defensible sampling design, and the planner places that calculation in the hands of those who allocate ship time. For data collection, a survey can be designed before a vessel leaves port by combining a prior estimate of ρ with the precision a program requires, converting a fixed budget of ship time into a station count or identifying the effort needed to reach a target precision. For monitoring and management, the same inversion shows where existing programs are over- or under-sampled relative to their objectives. For policy and reporting, the planner documents the assumptions (ρ , ε and calibration) behind a chosen effort, so that sampling

requirements in monitoring plans, survey proposals and emerging eDNA reporting standards can be stated transparently and revisited as new data accumulate.

Sampling effort required for eDNA surveys should be matched to the intended inferential goal. When the objective is to estimate GP parameters such as the mean (μ_d) for deriving depth-specific or region-specific abundance indices typically used in stock assessment indices (Shelton et al., 2022), the autocorrelation (ρ) for applied ecological questions, like implementing protected area networks or monitoring programs (Laroche, 2024), or the variance parameter (α^2), random sampling designs that adequately cover the full spatial domain of the species spatial autocorrelation are likely to be sufficient (Rocha et al., 2021), as sampling effort appears to have relatively limited influence on parameter estimation (Figure 5; Appendix S1: Section S4). In contrast, when the objective shifts toward predicting spatial distributions - such as identifying hotspots, mapping species occurrence, or quantifying spatial niche breadth - the accuracy of these predictions is much more sensitive to both sampling effort and the underlying spatial autocorrelation of the species (Figure 4A). Prediction is the more demanding of the two because it requires both accurate covariance parameters (Eq. 1–2) and observed points close enough to the location being predicted (Eq. 3).

This distinction also clarifies how sampling design shapes inference. Estimation of the mean benefits primarily from coverage (sampling across the full domain) and secondarily from balance (avoiding over-sampling a limited region) (Rocha et al., 2021). By contrast, estimation of the length-scale parameter requires not only broad coverage but also a wide distribution of pairwise distances, including closely spaced stations that resolve rapid decorrelation and widely spaced stations that constrain the long-distance tail of the covariance function (Kaufman & Shaby, 2013).

In practice, the starting point is a clear statement of the inferential objective and the precision it demands, since this determines whether effort should prioritize spatial coverage or the resolution of fine-scale autocorrelation. This is complemented by a prior estimate of the target taxon's spatial autocorrelation, drawn from published geostatistical fits, a pilot survey, or the depth-related pattern reported here, which the planner uses (Eq. 6–7) to characterize the relationship between sampling density (E) and expected accuracy. From this relationship, a sampling density can be identified beyond which additional effort yields diminishing returns relative to the precision required. Because the length-scale parameter cannot be recovered from a uniformly thinned grid, the resulting design should also be assessed for its inclusion of both widely and closely spaced station pairs. In multispecies programmes, these considerations are most appropriately governed by the patchiest taxon the survey is intended to resolve.

Applied to existing programs this yields a concrete conclusion. The West Coast eDNA grid analyzed here (3.6 samples per 10,000 km² per depth layer) and the CalCOFI 75-station grid (3.4 samples per 10,000 km²) already sit close to the effort at which returns diminish for smoothly distributed species and for estimating depth-specific means, so programs whose primary product is an abundance index need not expand. Programs intended to map patchily distributed taxa, including several forage and mesopelagic species examined here, are under-sampled by a factor of roughly two to three, and that shortfall is spatial rather than analytical (it cannot be recovered by sequencing more deeply or filtering more water at stations already occupied).

For managers and survey designers the practical message is that effort should be set by the patchiness of the taxa a programme is required to resolve and by the decision the data must

support, and that both can be specified in advance rather than discovered afterwards. Nevertheless, across the scenarios considered here and the spatial domain used in this study, sampling effort of 4 - 9 samples per 10,000 km² per depth layer marked the point at which diminishing returns became apparent given species-specific spatial autocorrelation. We therefore recommend that eDNA surveys report estimated length-scales alongside detection or concentration results; as such estimates accumulate, taxon- and region-specific priors will make prospective design routine rather than bespoke, and recalibrated coefficients can be entered directly into the planner.

Acknowledgments

This material is based upon research supported by the Office of Naval Research under Award Number N00014-22-1-2719 and support from the National Marine Fisheries Service's Genomic Strategic Initiative. Any use of trade, firm, or product names is for descriptive purposes only and does not imply endorsement by the U.S. Government.

Author contributions

Gledis Guri conceived the study, developed the methodology and the survey planner, performed the analyses and wrote the manuscript; Andrew O. Shelton contributed to the methodology, funding acquisition, data interpretation and discussion; Ryan P. Kelly contributed to conceptualization, funding acquisition, data interpretation and discussion; Owen R. Liu contributed to data interpretation and discussion; Krista M. Nichols contributed to conceptualization, funding acquisition and discussion. All authors contributed critically to the drafts and gave final approval for publication.

500 **Conflict of interest statement**

501 The authors declare no conflicts of interest.

502 **References**

- 503 Baetscher, D. S., Omori, K. L., Goethel, D. R., Shelton, A. O., Berger, A. M., Ledger, K. J.,
504 Nichols, K. M., & Larson, W. A. (2025). The Pragmatic Sceptic: A Practical Approach
505 for Integrating Environmental DNA Into Marine Stock Assessment and Fisheries
506 Management. *Fish and Fisheries*, faf.70001. <https://doi.org/10.1111/faf.70001>
- 507 Bessey, C., Jarman, S. N., Berry, O., Olsen, Y. S., Bunce, M., Simpson, T., Power, M.,
508 McLaughlin, J., Edgar, G. J., & Keesing, J. (2020). Maximizing fish detection with
509 eDNA metabarcoding. *Environmental DNA*, 2(4), 493–504.
510 <https://doi.org/10.1002/edn3.74>
- 511 Brandão-Dias, P. F. P., Hallack, D. M. C., Snyder, E. D., Tank, J. L., Bolster, D., Volponi, S.,
512 Shogren, A. J., Lamberti, G. A., Bibby, K., & Egan, S. P. (2023). Particle size influences
513 decay rates of environmental DNA in aquatic systems. *Molecular Ecology Resources*,
514 23(4), 756–770. <https://doi.org/10.1111/1755-0998.13751>
- 515 Dormann, C. F., McPherson, J. M., Araújo, M. B., Bivand, R., Bolliger, J., Carl, G., Davies, R.
516 G., Hirzel, A., Jetz, W., Daniel Kissling, W., Kühn, I., Ohlemüller, R., Peres-Neto, P. R.,
517 Reineking, B., Schröder, B., Schurr, F. M., & Wilson, R. (2007). Methods to account for
518 spatial autocorrelation in the analysis of species distributional data: A review. *Ecography*,
519 30(5), 609–628. <https://doi.org/10.1111/j.2007.0906-7590.05171.x>
- 520 Fortin, M.-J., Drapeau, P., & Legendre, P. (1989). Spatial autocorrelation and sampling design in
521 plant ecology. *Vegetatio*, 83(1–2), 209–222. <https://doi.org/10.1007/BF00031693>

- Fukaya, K., Murakami, H., Yoon, S., Minami, K., Osada, Y., Yamamoto, S., Masuda, R., Kasai, A., Miyashita, K., Minamoto, T., & Kondoh, M. (2021). Estimating fish population abundance by integrating quantitative data on environmental DNA and hydrodynamic modelling. *Molecular Ecology*, 30(13), 3057–3067. <https://doi.org/10.1111/mec.15530>
- García-Machado, E., Normandeau, E., Côté, G., Caza-Allard, I., Babin, C., & Bernatchez, L. (2024). Evaluating Sampling Designs to Survey Fish Diversity in Lakes From Northern Temperate Zones. *Environmental DNA*, 6(6), e70015. <https://doi.org/10.1002/edn3.70015>
- Gloor, G. B., Macklaim, J. M., Pawlowsky-Glahn, V., & Egozcue, J. J. (2017). Microbiome Datasets Are Compositional: And This Is Not Optional. *Frontiers in Microbiology*, 8, 2224. <https://doi.org/10.3389/fmicb.2017.02224>
- Gold, Z., Sprague, J., Kushner, D. J., Marin, E. Z., & Barber, P. H. (2021). eDNA metabarcoding as a biomonitoring tool for marine protected areas. *PLoS ONE*, 16(2 February 2021), 1–19. <https://doi.org/10.1371/journal.pone.0238557>
- Guri, G. (2026). *Smart eDNA sampling strategies: Spatial autocorrelation can guide efficient survey design* [Data set]. Zenodo. <https://doi.org/10.5281/zenodo.21141367>
- Guri, G., Liu, O. R., Kelly, R. P., Shaffer, M. R., Parsons, K. M., Ramón-Laca, A., Nichols, K. M., Brandão-Dias, P. F. P., Wells, A., & Shelton, A. O. (2025). *Quantitative, Multispecies Monitoring at a Continental Scale*. Ecology. <https://doi.org/10.1101/2025.10.30.685180>
- Guri, G., Shelton, A. O., Kelly, R. P., Yoccoz, N., Johansen, T., Præbel, K., Hanebrekke, T., Ray, J. L., Fall, J., & Westgaard, J.-I. (2024). Predicting trawl catches using environmental DNA. *ICES Journal of Marine Science*, fsae097. <https://doi.org/10.1093/icesjms/fsae097>

545 Jo, T. S., Murakami, H., & Nakadai, R. (2025). Spatial dispersal of environmental DNA particles
 546 in lentic and marine ecosystems: An overview and synthesis. *Ecological Indicators*, 174,
 547 113469. <https://doi.org/10.1016/j.ecolind.2025.113469>

548 Kaufman, C. G., & Shaby, B. A. (2013). The role of the range parameter for estimation and
 549 prediction in geostatistics. *Biometrika*, 100(2), 473–484.
 550 <https://doi.org/10.1093/biomet/ass079>

551 Kelly, R. P., Shelton, A. O., & Gallego, R. (2019). Understanding PCR Processes to Draw
 552 Meaningful Conclusions from Environmental DNA Studies. *Scientific Reports*, 9(1),
 553 1–14. <https://doi.org/10.1038/s41598-019-48546-x>

554 Kirtane, A., Wiczorek, D., Noji, T., Baskin, L., Ober, C., Plosica, R., Chenoweth, A., Lynch,
 555 K., & Sassoubre, L. (2021). Quantification of Environmental DNA (eDNA) shedding and
 556 decay rates for three commercially harvested fish species and comparison between eDNA
 557 detection and trawl catches. *Environmental DNA*, 3(6), 1142–1155.
 558 <https://doi.org/10.1002/edn3.236>

559 Laroche, F. (2024). Efficient sampling designs to assess biodiversity spatial autocorrelation:
 560 Should we go fractal? *Peer Community Journal*, 4, e76.
 561 <https://doi.org/10.24072/pcjournal.454>

562 Lewy, P., & Kristensen, K. (2009). Modelling the distribution of fish accounting for spatial
 563 correlation and overdispersion. *Canadian Journal of Fisheries and Aquatic Sciences*,
 564 66(10), 1809–1820. <https://doi.org/10.1139/F09-114>

565 Liu, O. R., Shelton, A. O., Ramón-Laca, A., Nichols, K. M., Ward, E. J., Phillips, E. M., Zamon,
 566 J. E., Wells, A., & Kelly, R. P. (2026). Mapping the marine distribution of eulachon

567 (Thaleichthys pacificus) in the Northeast Pacific using environmental DNA.
 568 *Communications Biology*, 9(1), 465. <https://doi.org/10.1038/s42003-026-09733-5>
 569 Paciorek, C. J. (2010). The Importance of Scale for Spatial-Confounding Bias and Precision of
 570 Spatial Regression Estimators. *Statistical Science*, 25(1). [https://doi.org/10.1214/10-](https://doi.org/10.1214/10-STS326)
 571 STS326
 572 Paciorek, C. J., & Schervish, M. J. (2006). Spatial modelling using a new class of nonstationary
 573 covariance functions. *Environmetrics*, 17(5), 483–506. <https://doi.org/10.1002/env.785>
 574 Perkins, N. R., Prall, M., Chakraborty, A., White, J. W., Baskett, M. L., & Morgan, S. G. (2021).
 575 Quantifying the statistical power of monitoring programs for marine protected areas.
 576 *Ecological Applications*, 31(1), e2215. <https://doi.org/10.1002/eap.2215>
 577 Perret, J., Charpentier, A., Pradel, R., Papuga, G., & Besnard, A. (2022). Spatially balanced
 578 sampling methods are always more precise than random ones for estimating the size of
 579 aggregated populations. *Methods in Ecology and Evolution*, 13(12), 2743–2756.
 580 <https://doi.org/10.1111/2041-210X.14015>
 581 R Core Team. (2024). *R: A language and environment for statistical computing*. R Foundation
 582 for Statistical Computing, Vienna, Austria. <https://www.R-project.org/>
 583 Radinger, J., Britton, J. R., Carlson, S. M., Magurran, A. E., Alcaraz-Hernández, J. D.,
 584 Almodóvar, A., Benejam, L., Fernández-Delgado, C., Nicola, G. G., Oliva-Paterna, F. J.,
 585 Torralva, M., & García-Berthou, E. (2019). Effective monitoring of freshwater fish. *Fish*
 586 *and Fisheries*, 20(4), 729–747. <https://doi.org/10.1111/faf.12373>
 587 Redfern, J., Ferguson, M., Becker, E., Hyrenbach, K., Good, C., Barlow, J., Kaschner, K.,
 588 Baumgartner, M., Forney, K., Ballance, L., Fauchald, P., Halpin, P., Hamazaki, T.,
 589 Pershing, A., Qian, S., Read, A., Reilly, S., Torres, L., & Werne, F. (2006). Techniques

for cetacean habitat modeling. *Marine Ecology Progress Series*, 310, 271–295.
<https://doi.org/10.3354/meps310271>

Rocha, A. D., Groen, T. A., Skidmore, A. K., & Willemsen, L. (2021). Role of Sampling Design When Predicting Spatially Dependent Ecological Data With Remote Sensing. *IEEE Transactions on Geoscience and Remote Sensing*, 59(1), 663–674.
<https://doi.org/10.1109/TGRS.2020.2989216>

Salter, I., Brandner, M., Djurhuus, A., Præbel, K., & Wangensteen, O. (2026). Comparison of Environmental DNA Metabarcoding and Trawl Data Highlights Gear Selectivity of Marine Demersal Fish Surveys on the Faroese Plateau. *Environmental DNA*, 8(2), e70253. <https://doi.org/10.1002/edn3.70253>

Schweiger, A. H., Irl, S. D. H., Steinbauer, M. J., Dengler, J., & Beierkuhnlein, C. (2016). Optimizing sampling approaches along ecological gradients. *Methods in Ecology and Evolution*, 7(4), 463–471. <https://doi.org/10.1111/2041-210X.12495>

Shaffer, M. R., Andruszkiewicz Allan, E., Van Cise, A. M., Parsons, K. M., Shelton, A. O., & Kelly, R. P. (2025). Observation Bias in Metabarcoding. *Molecular Ecology Resources*, 25(7), e14119. <https://doi.org/10.1111/1755-0998.14119>

Shelton, A. O., Ramón-Laca, A., Wells, A., Clemons, J., Chu, D., Feist, B. E., Kelly, R. P., Parker-Stetter, S. L., Thomas, R., Nichols, K. M., & Park, L. (2022). Environmental DNA provides quantitative estimates of Pacific hake abundance and distribution in the open ocean. *Proceedings of the Royal Society B: Biological Sciences*, 289(1971), 20212613. <https://doi.org/https://doi.org/10.1098/rspb.2021.2613>

Stan Development Team. (2023). *Stan Modeling Language Users Guide and Reference Manual* (Version 2.26.1) [Computer software]. <https://mc-stan.org/>

- Stein, A., & Ettema, C. (2003). An overview of spatial sampling procedures and experimental design of spatial studies for ecosystem comparisons. *Agriculture, Ecosystems & Environment*, 94(1), 31–47. [https://doi.org/10.1016/S0167-8809\(02\)00013-0](https://doi.org/10.1016/S0167-8809(02)00013-0)
- Taberlet, P., Coissac, E., Hajibabaei, M., & Rieseberg, L. H. (2012). Environmental DNA. *Molecular Ecology*, 21(8), 1789–1793. <https://doi.org/10.1111/j.1365-294X.2012.05542.x>
- Thorson, J. T., Shelton, A. O., Ward, E. J., & Skaug, H. J. (2015). Geostatistical delta-generalized linear mixed models improve precision for estimated abundance indices for West Coast groundfishes. *ICES Journal of Marine Science*, 72(5), 1297–1310. <https://doi.org/10.1093/icesjms/fsu243>
- Wagner, H. H., & Fortin, M.-J. (2005). Spatial analysis of landscapes: Concepts and statistics. *Ecology*, 86(8), 1975–1987. <https://doi.org/10.1890/04-0914>
- Xu, B., Zhang, C., Xue, Y., Ren, Y., & Chen, Y. (2015). Optimization of sampling effort for a fishery-independent survey with multiple goals. *Environmental Monitoring and Assessment*, 187(5). <https://doi.org/10.1007/s10661-015-4483-9>
- Yoccoz, N. G., Nichols, J. D., & Boulinier, T. (2001). Monitoring of biological diversity in space and time. *Trends in Ecology & Evolution*, 16(8), 446–453. [https://doi.org/10.1016/S0169-5347\(01\)02205-4](https://doi.org/10.1016/S0169-5347(01)02205-4)
- Zhao, J., Yang, K., & Ma, J. (2022). Optimization of Sampling Effort for Different Fishery Groups in the Yangtze River Estuary, China. *Marine and Coastal Fisheries*, 14(4). <https://doi.org/10.1002/mcf2.10214>

Figure captions

Figure 1. Conceptual illustration of how sampling density affects the reconstruction of (A) a patchy and (B) a smooth concentration field along a one-dimensional transect (km). Top rows show the true concentration (black line); middle and bottom rows show reconstructions (red dashed line) from few and many samples (green circles), respectively.


Figure 2. Workflow of the method for simulated (A) and empirical (B) data. (A) A field generated from known parameters on a grid (Q , $N = 18,916$) is sampled (Y , $N = 4,000$) to give a reference prediction (Z); ε_2 compares Z with predictions from thinned designs (Y_E), and ε_3 uses the prediction from $N = 358$ samples ($S_{E=358}$) as the reference. (B) The same thinning applied to the observed data ($N = 358$) quantifies ε_1 .

Figure 3. Gaussian process parameter estimates for the 12 species. (A) Spatial autocorrelation length-scale (ρ , km) by depth association (surface, midwater, deepwater). (B) Depth-specific mean eDNA concentration (μ ; log copies/L). (C) Amplitude (α^2) versus non-spatial variance (σ^2). Points are posterior means and bars 95% credible intervals; colors identify species.

Figure 4. (A) Prediction error (ε_2) across sampling densities (E , samples per 10,000 km² per depth layer) for simulated fields with different length-scales (colors); points are simulation outcomes and lines mean trends. (B) Marginal change in error ($\Delta\varepsilon_2$) per unit increase in E , predicted from Eq. 6 (Appendix S1: Figure S5); the dashed line marks the threshold ($\Delta\varepsilon_2/\Delta E = -0.05$) for negligible improvement.


Figure 5. Recovery of Gaussian process parameters in the simulations: mean root mean squared error (τ) between estimated and true values across 20 replicates for (A) ρ , (B) σ^2 , (C) depth-specific mean (μ_d) and (D) α^2 , by sampling density (E) and length-scale (colors).

657 **Figure 1**




658

659 **Figure 2**




661 **Figure 3**




662

663 **Figure 4**



664

665 **Figure 5**



666